



Integrating Different Samples in Single-cell RNA-seq

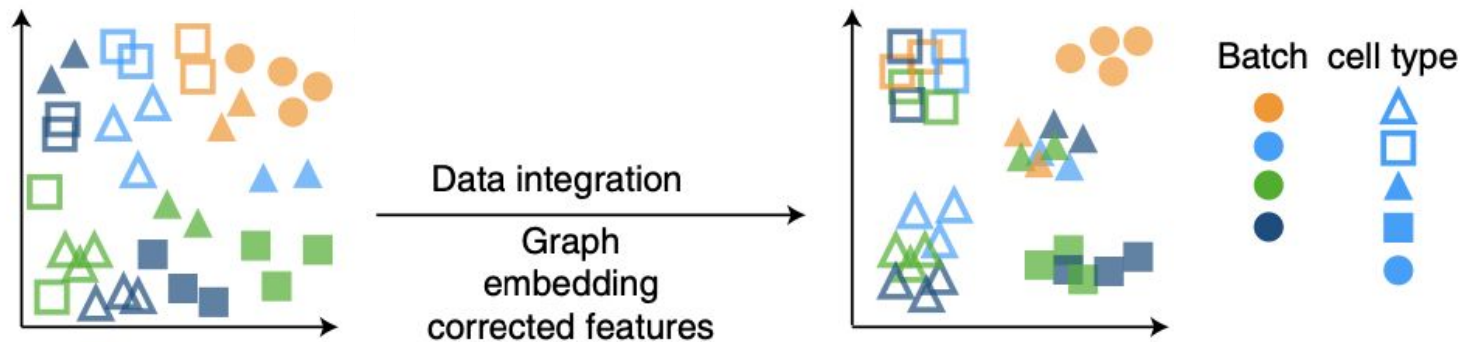
The Data Lab

Why integrate ("batch correct") datasets?

- Analyze multiple libraries at once rather than analyze each library individually
- Integration allows us to:
 - Perform clustering on a single integrated dataset
 - Visualize integrated libraries in one space
 - Perform downstream analyses that would benefit from a larger sample size:
 - Identify marker genes
 - Cell type annotation
- Provides an integrated profile of multiple libraries, which is often represented as a reduced dimension matrix of some kind (e.g., PCA or "latent embedding")
- Note that we can expect *tradeoffs* between reducing technical variation and retaining biological variation


What does "good integration" look like?

- Helps mitigate (👉) the *batch effects* caused by variation across libraries, while hopefully (👉) still preserving biological information
 - Technical variation can arise from *all that comes with* separate library preps and sequencing



Adapted from Lueken et al., 2022

Example of (what looks like*) successful integration



Gayoso *et al.*, 2021

*Measuring success is actually kind of tricky! Stay tuned...

Let's have a closer look at methods we'll be using

- **MNN: Mutual nearest neighbors**



- Specifically, we'll use FastMNN 🚀
- Haghverdi, L, Lun, A, Morgan, M, et al. *Batch effects in single-cell RNA-sequencing data are corrected by matching mutual nearest neighbors.* (2018) <https://doi.org/10.1038/nbt.4091>

- **Harmony**

- Korsunsky, I, Millard, N, Fan, J, et al. *Fast, sensitive and accurate integration of single-cell data with Harmony.* (2019) <https://doi.org/10.1038/s41592-019-0619-0>



Mutual nearest neighbors batch correction

- Imagine we have 2 batches, each with 3 cell types
 - Red and blue are shared but yellow and pink are not!
 - Before beginning integration, cosine distances are first calculated among pairs of cells *within each sample*
 - This enables expression profile comparisons and sets up the data for integration
- First, we identify pairs of cells with mutually similar expression profiles
 - These are our "mutual nearest neighbors"




Mutual nearest neighbors batch correction

- Next, compute a batch correction vector for each MNN pair
- Finally, calculate the weighted average of these vectors to get cell-specific batch corrections to perform the final integration
 - Note that **w** and **z** don't "look" as "integrated"! Why?



Some assumptions that MNN makes

- At least one cell population is present in both batches
- The batch effect is almost orthogonal to the biological effects
 - Roughly means, batches and biology are expected to have *separate variation*
- The batch-effect variation is much smaller than the biological-effect variation across cell types



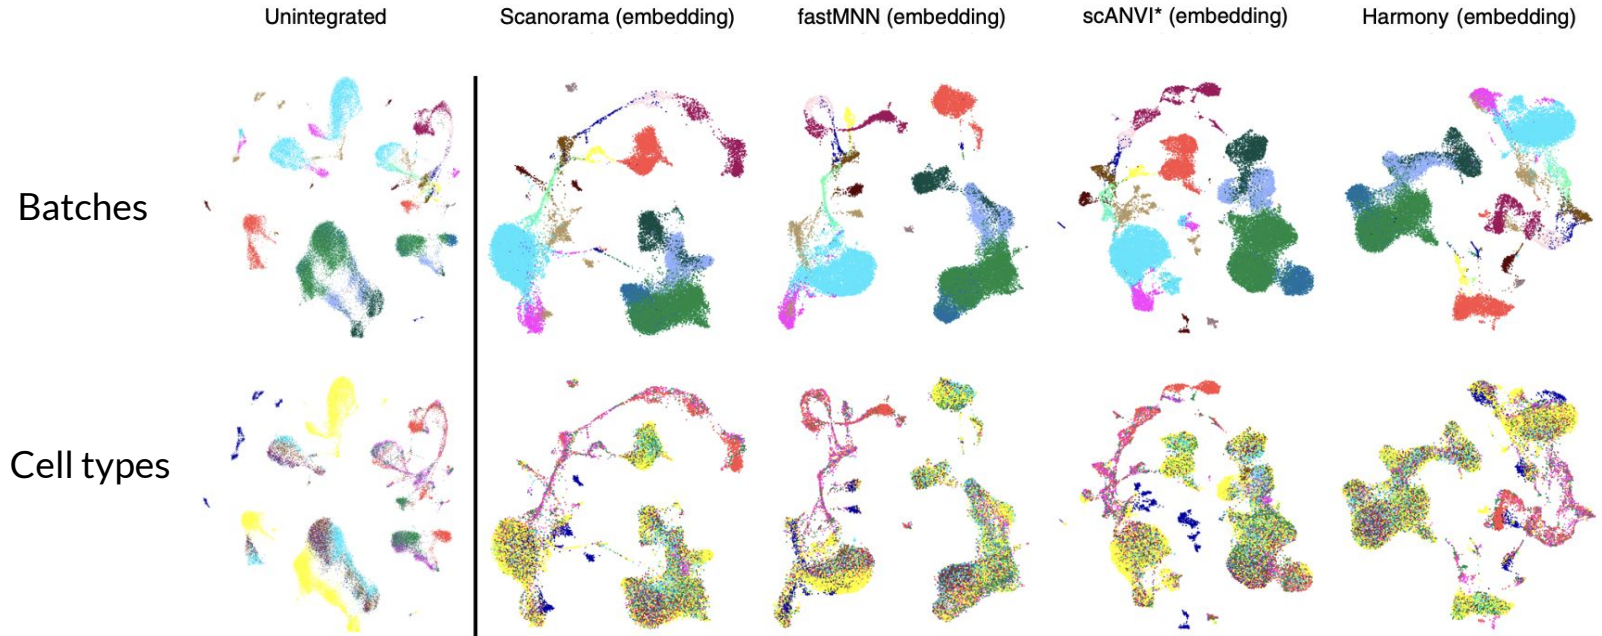
Evaluating integration: What counts as "good"?

- Before and after UMAP vibes



- We expect that batches cluster less after integration due to removed technical variation
- But "biology" across batches should still cluster together
 - Biology = cell types (which may not be known!), tissue type, donor, etc.

Your mileage may vary across methods!








Luecken, M.D., Büttner, M., Chaichoompu, K. et al. *Benchmarking atlas-level data integration in single-cell genomics*. (2022). <https://doi.org/10.1038/s41592-019-0619-0>

We performed some benchmarking on simulated data from Luecken *et al.*

- We evaluated several methods, four of which we'll show here:
 - **FastMNN**
 - **Harmony**
 - **Seurat using CCA** (canonical correlation analysis)
 - **Seurat using RPCA** (reciprocal PCA)
 - (We'll note that we also looked at scVI, which seemed to work well but is slow on CPU and it's in Python which is beyond the scope of our workshop!)
- We chose these methods based on performance in Luecken *et al.* and their usability






Scenario 1: All cell types are present in all batches

UMAPs colored by Batch








Scenario 1: All cell types are present in all batches

UMAPs colored by Cell Type




Scenario 2: Cell types are not present in all batches, and not all batches have cells in common

UMAPs colored by Batch



Scenario 2: Cell types are not present in all batches, and not all batches have cells in common

UMAPs colored by Cell Type



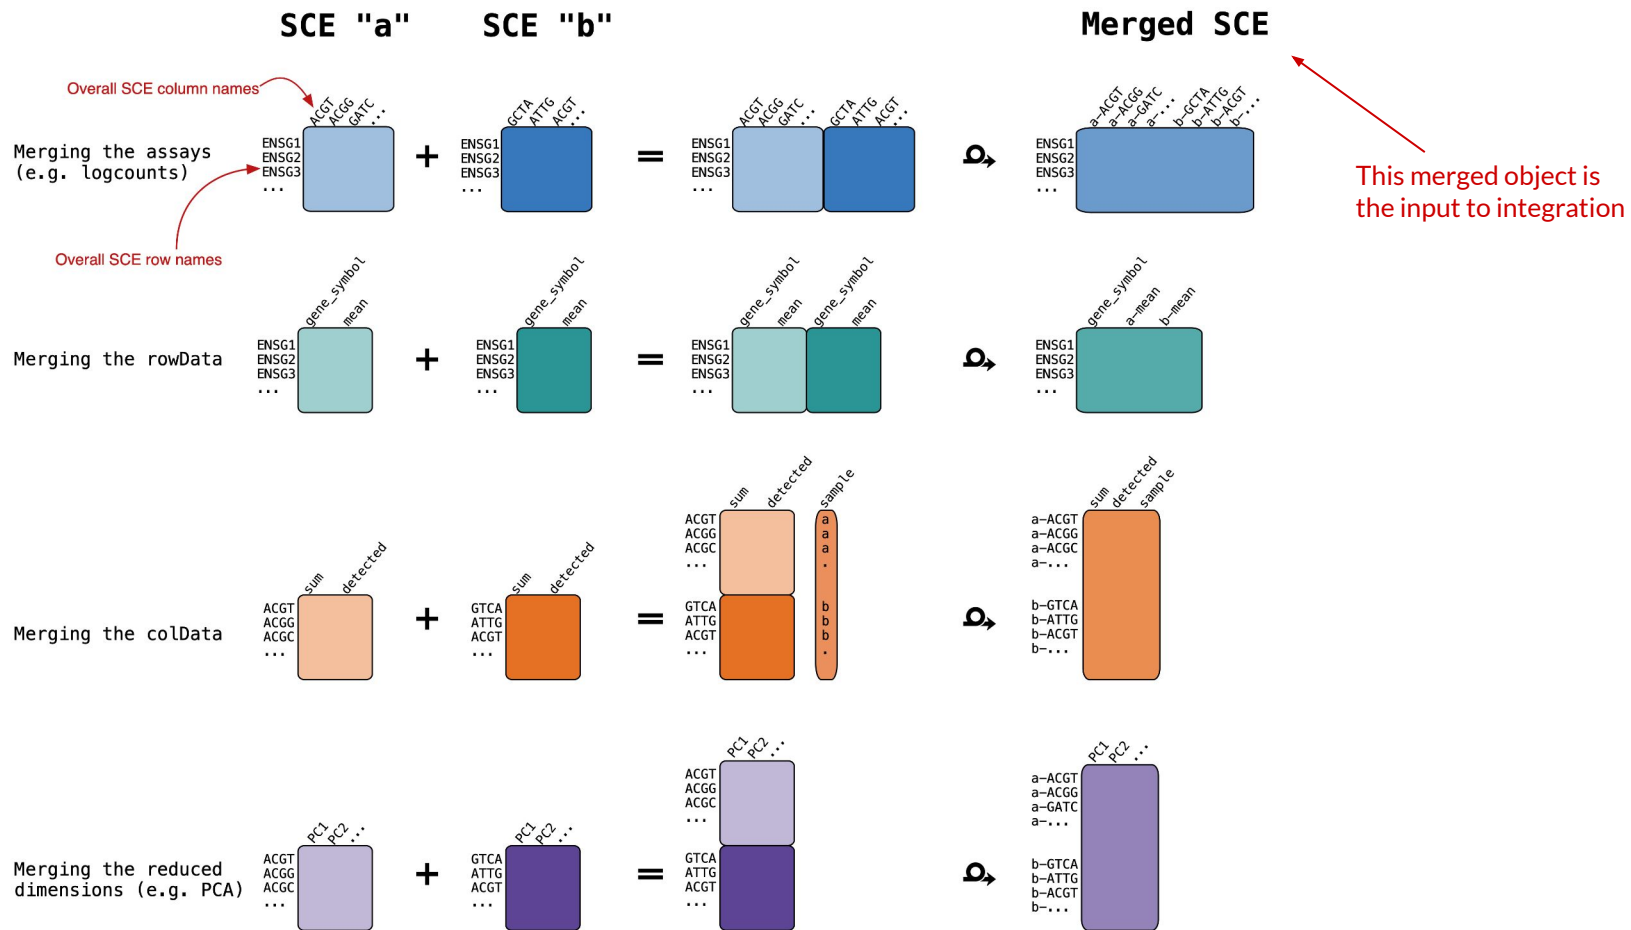
Evaluating integration: What counts as "good"?



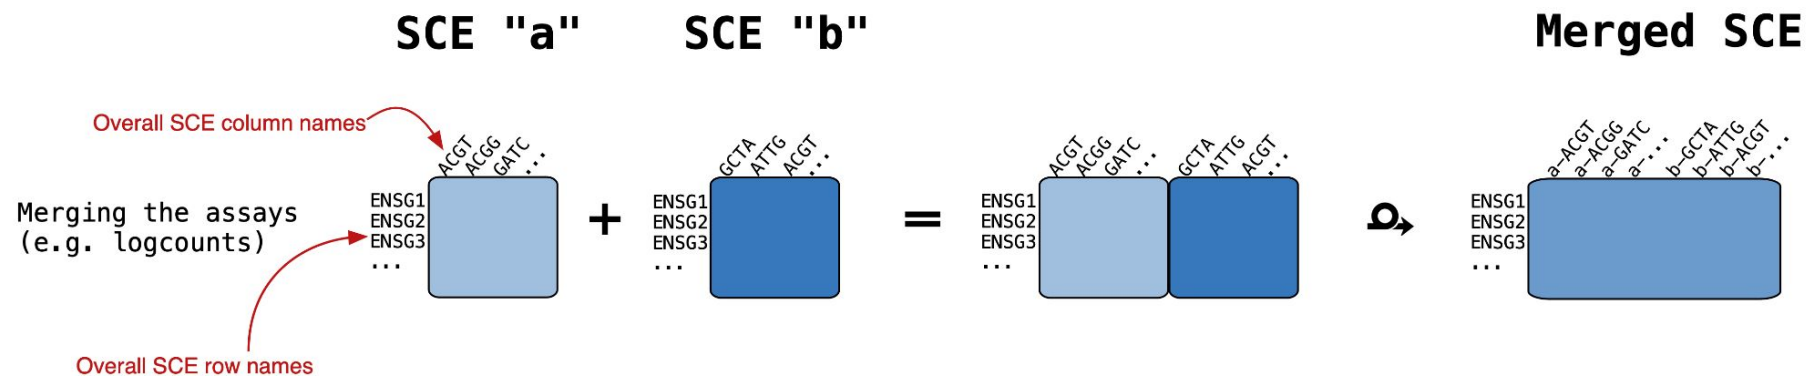
- Before and after UMAP vibes
 - We expect that batches cluster less after integration due to removed technical variation
 - But "biology" across batches should still cluster together
 - Biology = cell types (which may not be known!), tissue type, donor, etc.
- There are several metrics for evaluating batch correction
 - Luecken *et al.* (2022) is an excellent reference <https://doi.org/10.1038/s41592-019-0619-0>
 - Caution: Metrics generally do not measure "was integration successful," but other proxies which *sometimes can help us tell* if integration was successful (or at least not unsuccessful)

Performing integration: Bookkeeping


- As input, many methods (in R!) require you to merge all SCEs into one *unintegrated SCE object*, which can then be integrated
 - Key point: Combining is NOT integrating
- This means SCEs need to be able to be merged, which may require us to manipulate SCE objects first for compatibility!



Merging SCE assays




Per-gene (feature) data: Each row is a gene



Per-cell data: Each row is a cell

Merging the colData



Merging the reduced dimensions (e.g. PCA)

